

Leo (Yi-Ting) Chao

Irvine, CA | +1-858-519-4030 | leo.yt.chao@gmail.com | [linkedin.com/in/leoytchao](https://www.linkedin.com/in/leoytchao) | twleochao.github.io

Education

University of California, San Diego | La Jolla, CA Sep 2026 - Present

Master of Science, Computer Science and Engineering

University of California, Irvine | Irvine, CA Sep 2023 - Jun 2026

Bachelor of Science, Computer Science - Intelligent Systems (GPA: 3.96/4.00) | Dean's Honors List, Accelerated 3-Year Degree

Skills

Programming Languages: Python, C++

LLM/Agents: Google ADK, Vertex AI, PyTorch, HuggingFace Transformers, LangGraph

Interpretability/Eval: TransformerLens, SAELens, LLM-as-a-Judge pipelines

Tools: GCP, Docker, Git, Linux, FastAPI

Research Experience

ARCALA Research Lab - UC Irvine | Irvine, CA *Undergraduate Research Assistant* Jan 2025 - Present

- Co-developed Chat-Bloom, an LLM evaluation taxonomy derived from Bloom's taxonomy, to classify and measure behavioral shifts in human-LLM interaction logs; organized and analyzed a dataset of 8,076 prompt-response pairs
- Trained BERT/DistilBERT ensemble classifiers for automated interaction taxonomy labeling, achieving ~90% accuracy validated against human-coder baseline (75.9% inter-rater reliability)
- Built PREVAIL, a predictive-replay harness to measure next-turn behavioral shift in LLM-assisted workflows
- Manuscript: Chao, L., et al.** "Chat-Bloom: A Taxonomy for Classifying Cognitive Offloading in Student-AI interactions in CS Education." (Revision in progress; target resubmission May 2026)

Professional Experience

Moovv | Irvine, CA, *AI Intern* Jun - Aug 2026 (Incoming)

Acer | Taipei, Taiwan, *AI Systems Intern* Jun - Aug 2025

- Architected a multi-agent SDLC automation system using Google ADK and Vertex AI, reducing manual iteration time by 35%; deployed to GCP with a production path into internal Google Agentspace
- Refactored LangGraph prototype into production-grade ADK workflow with persistent memory management, integrating 6 enterprise GCP services including Cloud Storage and Firestore
- Designed state management architecture for long-horizon agent tasks, ensuring context preservation across 10+ dynamic subtask agents executing code generation, API stub creation, and test drafting

EZ Intelligent Technology | Taipei, Taiwan, *Software Engineering Intern* Jun - Aug 2024

- Co-developed a real-time pedestrian flow management system deployed at Keelung Harbor using PyTorch and OpenCV, achieving 98% validation accuracy
- Boosted in-motion detection generalization by 20% through dropout tuning and image preprocessing pipeline refinement

Technical Projects

LLM Activation Steering | Mechanistic Interpretability (*Python, TransformerLens, SAELens, PyTorch*) Jan - Mar 2026

- Built an inference-time intervention pipeline to inject SAE feature vectors into Gemma's residual stream, evaluating sentiment and truthfulness steering across 1,372 IMDb and 817 TruthfulQA samples
- Designed an LLM-as-a-Judge evaluation framework using GPT-5 with multi-dimensional metrics: relative sentiment shift, atomic truthfulness verification, and normalized coherence scoring

HeadsUp | Real-Time Poker Hud (*Python, XGBoost, PySide6*) Jul - Dec 2025

- Trained a GTO-Lite XGBoost model on a 50,000-sample dataset generated via a custom eval7-based oracle, achieving 95.6% test accuracy in postflop action classification
- Built a multi-threaded inference overlay delivering <20ms end-to-end latency by isolating game state polling, model inference, and UI rendering in separate threads